

GRADIENT BOOSTING FRAMEWORK

```
import warnings
warnings.filterwarnings('ignore')
```

Load dataset

```
import pandas as pd
from sklearn.datasets import load_breast_cancer, fetch_california_housing

# reg - dataset
california_housing = fetch_california_housing(as_frame=True)
housing_df = california_housing.frame
print("REGRESSION")
display(housing_df.head())

# cls - dataset
breast_cancer = load_breast_cancer(as_frame=True)
Bcancer_df = breast_cancer.frame
print("\nCLASSIFICATION")
display(Bcancer_df.head())
```

REGRESSION

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MedHouseVal	
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422	

CLASSIFICATION

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	...	worst texture	worst perimeter	worst area
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871	...	17.33	184.60	2011.0
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667	...	23.41	158.80	1959.0
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999	...	25.53	152.50	1709.0
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.09744	...	26.50	98.87	568.0
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883	...	16.67	152.20	1575.0

5 rows × 31 columns

shape

```
print("reg", housing_df.shape)
print("cls", Bcancer_df.shape)

reg (20640, 9)
cls (569, 31)
```

info

```
housing_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
```

```
#   Column      Non-Null Count  Dtype
---  -
0   MedInc      20640 non-null   float64
1   HouseAge     20640 non-null   float64
2   AveRooms     20640 non-null   float64
3   AveBedrms    20640 non-null   float64
4   Population   20640 non-null   float64
5   AveOccup     20640 non-null   float64
6   Latitude     20640 non-null   float64
7   Longitude    20640 non-null   float64
8   MedHouseVal  20640 non-null   float64
dtypes: float64(9)
memory usage: 1.4 MB
```

```
Bcancer_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   mean radius  569 non-null   float64
1   mean texture 569 non-null   float64
2   mean perimeter 569 non-null   float64
3   mean area    569 non-null   float64
4   mean smoothness 569 non-null   float64
5   mean compactness 569 non-null   float64
6   mean concavity 569 non-null   float64
7   mean concave points 569 non-null   float64
8   mean symmetry 569 non-null   float64
9   mean fractal dimension 569 non-null   float64
10  radius error  569 non-null   float64
11  texture error 569 non-null   float64
12  perimeter error 569 non-null   float64
13  area error    569 non-null   float64
14  smoothness error 569 non-null   float64
15  compactness error 569 non-null   float64
16  concavity error 569 non-null   float64
17  concave points error 569 non-null   float64
18  symmetry error 569 non-null   float64
19  fractal dimension error 569 non-null   float64
20  worst radius  569 non-null   float64
21  worst texture 569 non-null   float64
22  worst perimeter 569 non-null   float64
23  worst area    569 non-null   float64
24  worst smoothness 569 non-null   float64
25  worst compactness 569 non-null   float64
26  worst concavity 569 non-null   float64
27  worst concave points 569 non-null   float64
28  worst symmetry 569 non-null   float64
29  worst fractal dimension 569 non-null   float64
30  target       569 non-null   int64
dtypes: float64(30), int64(1)
memory usage: 137.9 KB
```

```
housing_df["MedHouseVal"].median()
```

```
1.797
```

## ▼ remove noise columns

```
housing_df.drop(columns=["Population", "Latitude", "Longitude"],inplace=True)
```

```
Bcancer_df.drop(columns=[col for col in Bcancer_df.columns if "error" in col],inplace=True)
```

```
Bcancer_df.columns
```

```
Index(['mean radius', 'mean texture', 'mean perimeter', 'mean area',
      'mean smoothness', 'mean compactness', 'mean concavity',
      'mean concave points', 'mean symmetry', 'mean fractal dimension',
      'worst radius', 'worst texture', 'worst perimeter', 'worst area',
      'worst smoothness', 'worst compactness', 'worst concavity',
      'worst concave points', 'worst symmetry', 'worst fractal dimension',
      'target'],
      dtype='object')
```

```
print("reg",housing_df.shape)
print("cls",Bcancer_df.shape)
```


```
y_pred_c = xgb_clf.predict(X_test_bcancer)
print("XGBoost Classification Accuracy:", accuracy_score(y_test_bcancer, y_pred_c))
```

XGBoost Classification Accuracy: 0.956140350877193

```
plt.figure(figsize=(15,8))
plot_tree(xgb_clf, num_trees=0)
plt.title("XGBoost Classification Tree (Breast Cancer)")
plt.show()
```

<Figure size 1500x800 with 0 Axes>


XGBoost Classification Tree (Breast Cancer)



```
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test_bcancer, y_pred_c)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.title("XGBoost Classification Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

XGBoost Classification Confusion Matrix



## ✓ LightGBM

### REGRESSION

```
import lightgbm as lgb
from sklearn.metrics import mean_squared_error, r2_score

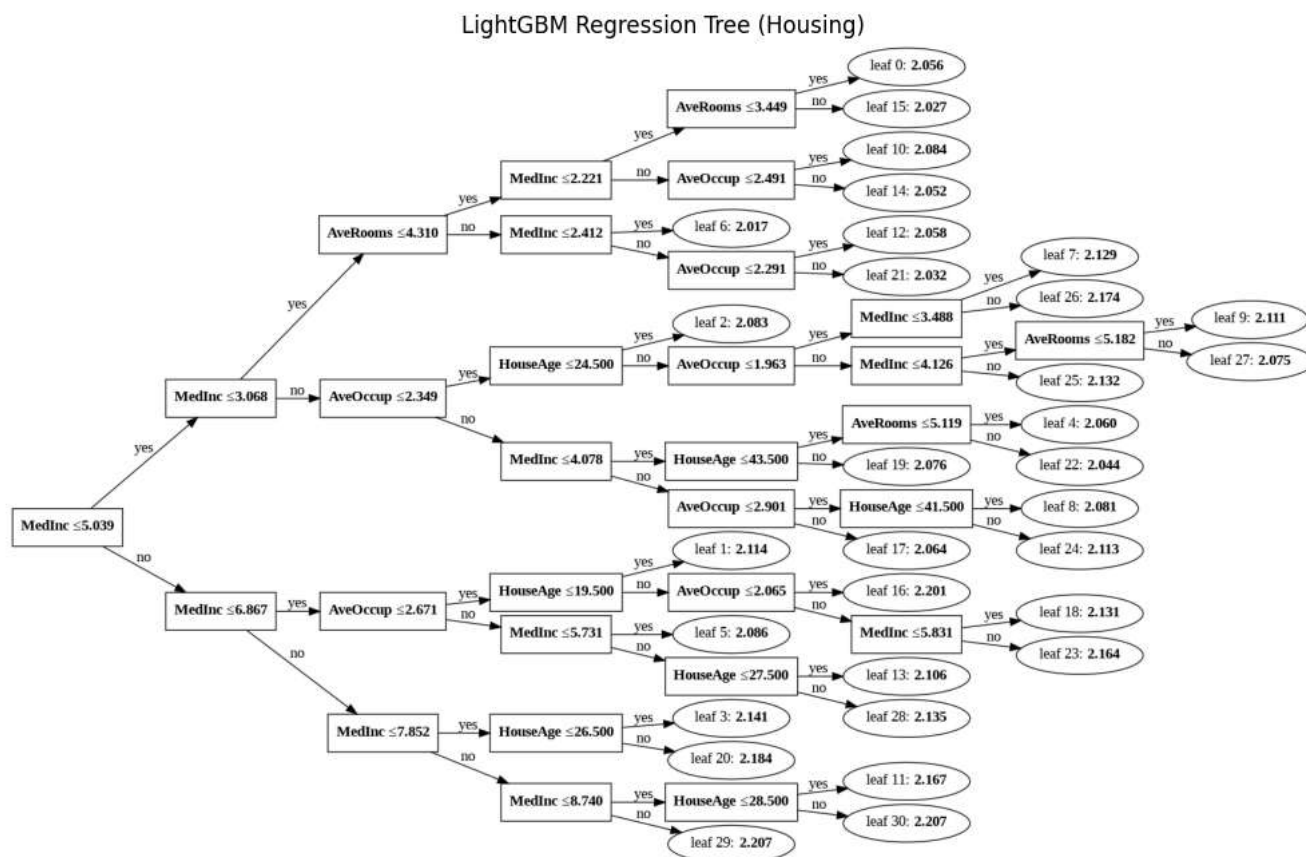
lgb_reg = lgb.LGBMRegressor(
    n_estimators=200,
    learning_rate=0.05,
    max_depth=-1,
    num_leaves=31,
    min_child_samples=5,
    random_state=42,
    verbose=-1)
```

```
lgb_reg.fit(X_train_housing, y_train_housing)
y_pred_L = lgb_reg.predict(X_test_housing)
print("LightGBM Regression MSE:", mean_squared_error(y_test_housing, y_pred_L))
print("R2:", r2_score(y_test_housing, y_pred_L))
```

LightGBM Regression MSE: 0.4189918915761474  
R2: 0.6802586837136873

```
import matplotlib.pyplot as plt
import lightgbm as lgb
from lightgbm import plot_tree

ax = plot_tree(lgb_reg, tree_index=0, figsize=(15,8))
plt.title("LightGBM Regression Tree (Housing)")
plt.show()
```



## CLASSIFICATION

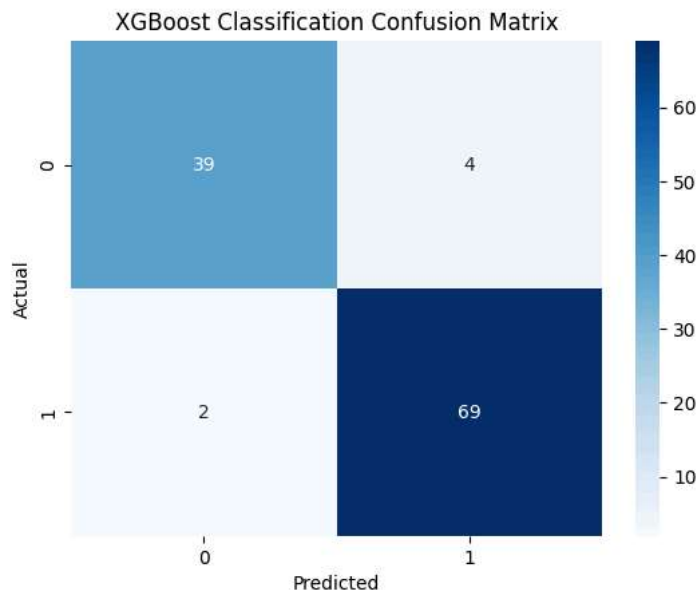
```
lgb_clf = lgb.LGBMClassifier(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)
lgb_clf.fit(X_train_bcancer, y_train_bcancer)

y_pred_LC = lgb_clf.predict(X_test_bcancer)
print("LightGBM Classification Accuracy:", accuracy_score(y_test_bcancer, y_pred_LC))
```

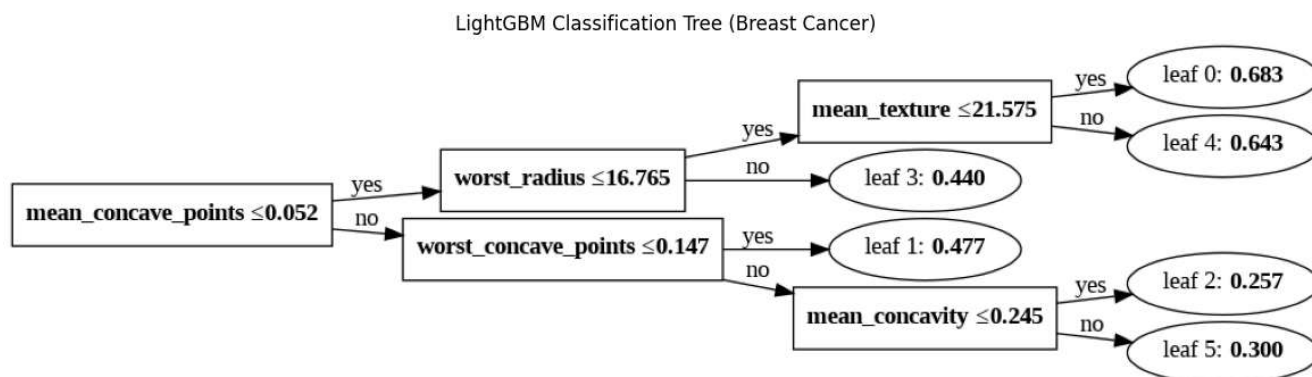
LightGBM Classification Accuracy: 0.9473684210526315

```
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test_bcancer, y_pred_LC)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.title("XGBoost Classification Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```
ax = plot_tree(lgb_clf, tree_index=0, figsize=(15,8))
plt.title("LightGBM Classification Tree (Breast Cancer)")
plt.show()
```



## ✓ CatBoost

```
pip install catboost
```

```

Requirement already satisfied: catboost in /usr/local/lib/python3.12/dist-packages (1.2.8)
Requirement already satisfied: graphviz in /usr/local/lib/python3.12/dist-packages (from catboost) (0.21)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (from catboost) (3.10.0)
Requirement already satisfied: numpy<3.0,>=1.16.0 in /usr/local/lib/python3.12/dist-packages (from catboost) (2.0.2)
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.12/dist-packages (from catboost) (2.2.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from catboost) (1.16.2)
Requirement already satisfied: plotly in /usr/local/lib/python3.12/dist-packages (from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.12/dist-packages (from catboost) (1.17.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost) (2.9)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (4.60.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (25.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (3.2.4)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.12/dist-packages (from plotly->catboost) (8.5.0)
  
```

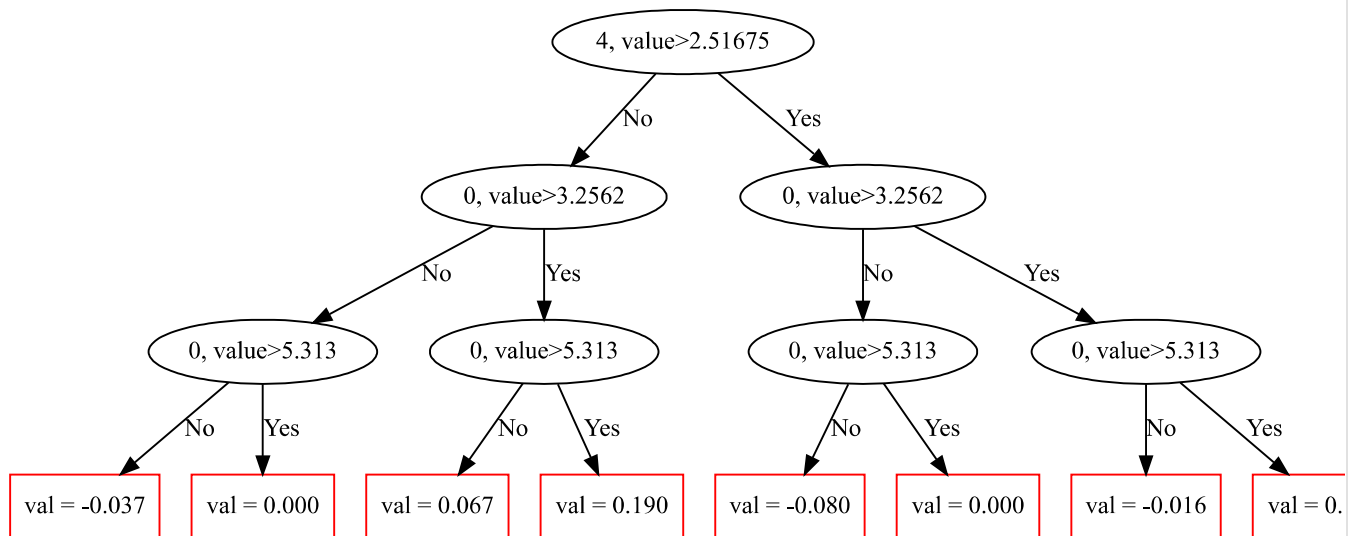
```
from catboost import CatBoostRegressor

cat_reg = CatBoostRegressor(iterations=100, learning_rate=0.1, depth=3, verbose=0, random_state=42)
cat_reg.fit(X_train_housing, y_train_housing)

y_pred_cat = cat_reg.predict(X_test_housing)
print("CatBoost Regression MSE:", mean_squared_error(y_test_housing, y_pred_cat))
print("R2:", r2_score(y_test_housing, y_pred_cat))
```

CatBoost Regression MSE: 0.4504888076209366  
R2: 0.656227403038132

```
cat_reg.plot_tree(tree_idx=0)
```



## CLASSIFICATION

```
from catboost import CatBoostClassifier

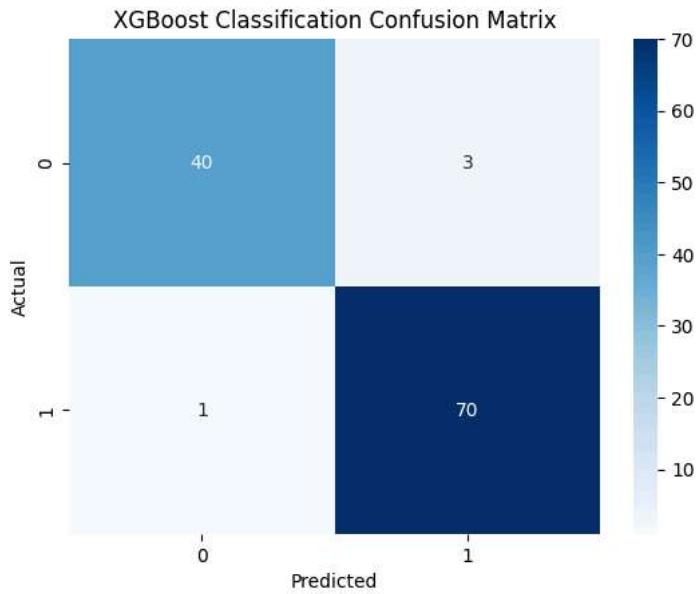
cat_clf = CatBoostClassifier(iterations=100, learning_rate=0.1, depth=3, verbose=0, random_state=42)
cat_clf.fit(X_train_bcancer, y_train_bcancer)

y_pred_catc = cat_clf.predict(X_test_bcancer)
print("CatBoost Classification Accuracy:", accuracy_score(y_test_bcancer, y_pred_catc))
```

CatBoost Classification Accuracy: 0.9649122807017544

```
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test_bcancer, y_pred_catc)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.title("XGBoost Classification Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```
cat_clf.plot_tree(tree_idx=0)
```

